

Aprendizado Profundo

Frameworks de Desenvolvimento de Redes Neurais - PyTorch

Professor: Lucas Silveira Kupssinskü

Agenda

- Redes Convolucionais
 - Camadas Convolucionais
 - Camadas de Pooling
- De Jupyter Notebook para *script* de treinamento
- Utilitários monitorar o *hardware* durante o experimento

Redes Convolucionais

- São redes que incluem operações de convolução
 - Usadas principalmente em imagens
- Criadas em 1989(!) por Yann LeCun (LeNet - 5)
 - LECUN, Yann et al. Handwritten digit recognition with a back-propagation network. Advances in neural information processing systems, v. 2, 1989.
- Explodiram em popularidade pós 2012
 - AlexNet
 - KRIZHEVSKY, Alex; SUTSKEVER, Ilya; HINTON, Geoffrey E. Imagenet classification with deep convolutional neural networks. Advances in neural information processing systems, v. 25, 2012.
 - Treinamento em GPUs
 - +Dados
 - +Poder Computacional

Redes Convolucionais

- Antigas conhecidas nossas
 - Tamanho do Kernel
 - Número de filtros
 - *Stride*
 - *Padding*
 - *Dilation*

| | | |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

Camadas Convolucionais

- Temos diversas camadas convolucionais disponíveis no PyTorch

Convolution Layers

`nn.Conv1d`

Applies a 1D convolution over an input signal composed of several input planes.

`nn.Conv2d`

Applies a 2D convolution over an input signal composed of several input planes.

`nn.Conv3d`

Applies a 3D convolution over an input signal composed of several input planes.

`nn.ConvTranspose1d`

Applies a 1D transposed convolution operator over an input image composed of several input planes.

`nn.ConvTranspose2d`

Applies a 2D transposed convolution operator over an input image composed of several input planes.

`nn.ConvTranspose3d`

Applies a 3D transposed convolution operator over an input image composed of several input planes.

Camadas Convolucionais

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,  
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

where \star is the valid 2D **cross-correlation** operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

Camadas de Pooling

`nn.MaxPool1d`

Applies a 1D max pooling over an input signal composed of several input planes.

`nn.MaxPool2d`

Applies a 2D max pooling over an input signal composed of several input planes.

`nn.MaxPool3d`

Applies a 3D max pooling over an input signal composed of several input planes.

`nn.MaxUnpool1d`

Computes a partial inverse of `MaxPool1d`.

`nn.MaxUnpool2d`

Computes a partial inverse of `MaxPool2d`.

`nn.MaxUnpool3d`

Computes a partial inverse of `MaxPool3d`.

`nn.AvgPool1d`

Applies a 1D average pooling over an input signal composed of several input planes.

`nn.AvgPool2d`

Applies a 2D average pooling over an input signal composed of several input planes.

Camadas de Pooling

```
CLASS torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False,  
                        ceil_mode=False) [SOURCE]
```


Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C, H, W) , output (N, C, H_{out}, W_{out}) and `kernel_size` (kH, kW) can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} input(N_i, C_j, stride[0] \times h + m, stride[1] \times w + n)$$

If `padding` is non-zero, then the input is implicitly padded with negative infinity on both sides for `padding` number of points. `dilation` controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

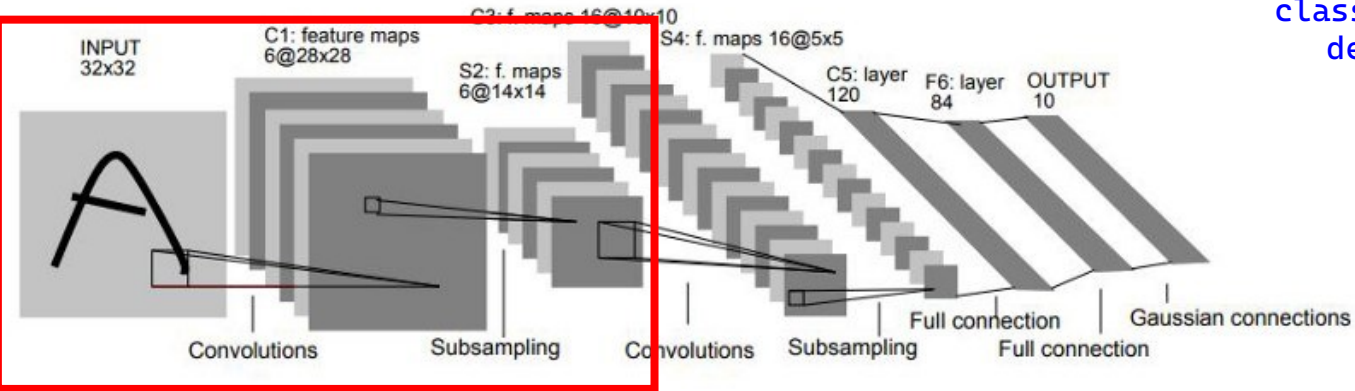
LeNet-5



```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)

    def forward(self, x):
        x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
        x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
        x = torch.flatten(x, 1)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```


LeNet-5



```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)
```

```
def forward(self, x):
    x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
    x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
    x = torch.flatten(x, 1)
    x = torch.relu(self.fc1(x))
    x = torch.relu(self.fc2(x))
    x = self.fc3(x)
    return x
```

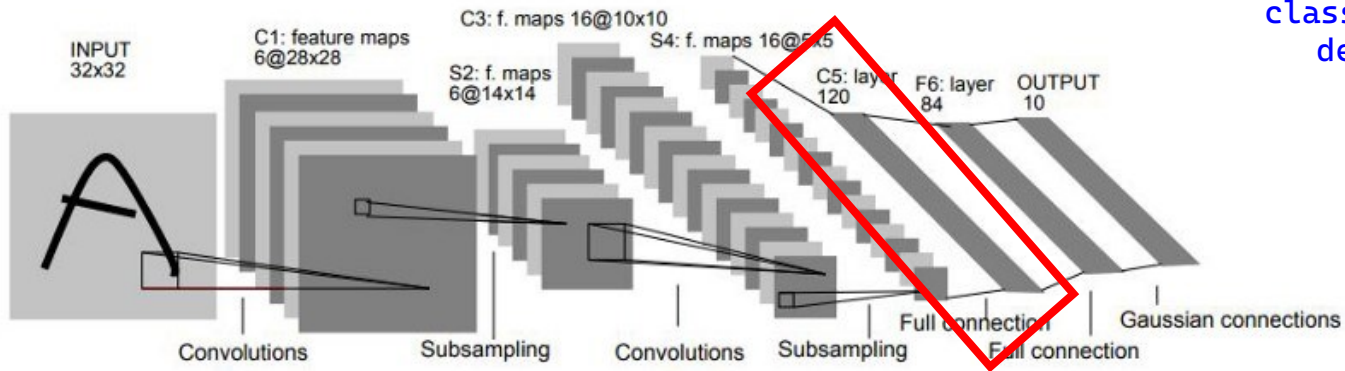
LeNet-5



```
class LeNet5(torch.nn.Module):  
    def __init__(self):  
        super(LeNet5, self).__init__()  
        self.conv1 = torch.nn.Conv2d(  
            in_channels=1, out_channels=6, kernel_size=5, padding=2)  
        self.conv2 = torch.nn.Conv2d(  
            in_channels=6, out_channels=16, kernel_size=5)  
        self.fc1 = torch.nn.Linear(16*5*5, 120)  
        self.fc2 = torch.nn.Linear(120, 84)  
        self.fc3 = torch.nn.Linear(84, 10)
```

```
    def forward(self, x):  
        x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)  
        x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)  
        x = torch.flatten(x, 1)  
        x = torch.relu(self.fc1(x))  
        x = torch.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x
```

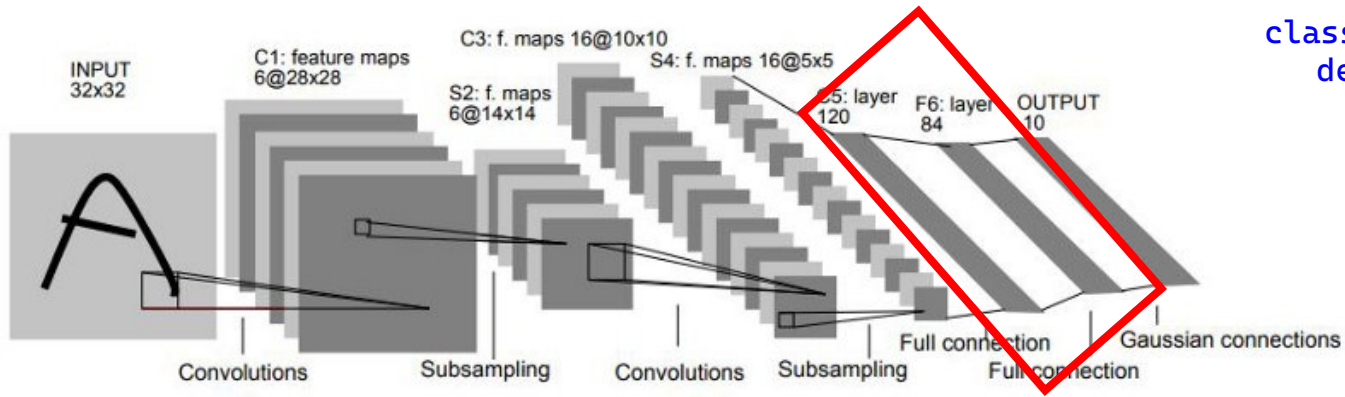
LeNet-5



```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)

    def forward(self, x):
        x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
        x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
        x = torch.flatten(x, 1)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```


LeNet-5



```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)
```

```
def forward(self, x):
    x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
    x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
    x = torch.flatten(x, 1)
    x = torch.relu(self.fc1(x))
    x = torch.relu(self.fc2(x))
    x = self.fc3(x)
    return x
```

LeNet-5



```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)
```

```
def forward(self, x):
    x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
    x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
    x = torch.flatten(x, 1)
    x = torch.relu(self.fc1(x))
    x = torch.relu(self.fc2(x))
    x = self.fc3(x)
    return x
```

LeNet-5

- Essa implementação difere da LeNet-5 original:
 - Função de ativação tanh
 - *Average Pooling*

```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = torch.nn.Conv2d(
            in_channels=1, out_channels=6, kernel_size=5, padding=2)
        self.conv2 = torch.nn.Conv2d(
            in_channels=6, out_channels=16, kernel_size=5)
        self.fc1 = torch.nn.Linear(16*5*5, 120)
        self.fc2 = torch.nn.Linear(120, 84)
        self.fc3 = torch.nn.Linear(84, 10)

    def forward(self, x):
        x = torch.max_pool2d(torch.relu(self.conv1(x)), 2)
        x = torch.max_pool2d(torch.relu(self.conv2(x)), 2)
        x = torch.flatten(x, 1)
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

LeNet-5

- Podemos encapsular partes da rede em um `nn.Sequential`

```
class LeNet5(torch.nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.feature_extractor = torch.nn.Sequential(
            torch.nn.Conv2d(
                in_channels=1, out_channels=6, kernel_size=5, padding=2),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2),
            torch.nn.Conv2d(
                in_channels=6, out_channels=16, kernel_size=5),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2)
        )
        self.classifier = torch.nn.Sequential(
            torch.nn.Flatten(1),
            torch.nn.Linear(16*5*5, 120),
            torch.nn.ReLU(),
            torch.nn.Linear(120, 84),
            torch.nn.ReLU(),
            torch.nn.Linear(84, 10)
        )

    def forward(self, x):
        x = self.feature_extractor(x)
        x = self.classifier(x)
        return x
```


MNIST

- Vamos usar a `torchvision.datasets` para carregar os dados
- Essa classe facilita o download de alguns conjuntos de dados bastante utilizados em problemas de visão computacional

```
CLASS torchvision.datasets.MNIST(root: str, train: bool = True, transform: Union[Callable,
NoneType] = None, target_transform: Union[Callable, NoneType] = None, download: bool [SOURCE]
= False) → None
```

MNIST Dataset.

Parameters:

- **root** (*string*) – Root directory of dataset where `MNIST/processed/training.pt` and `MNIST/processed/test.pt` exist.
- **train** (*bool, optional*) – If True, creates dataset from `training.pt`, otherwise from `test.pt`.
- **download** (*bool, optional*) – If true, downloads the dataset from the internet and puts it in root directory. If dataset is already downloaded, it is not downloaded again.
- **transform** (*callable, optional*) – A function/transform that takes in an PIL image and returns a transformed version. E.g, `transforms.RandomCrop`
- **target_transform** (*callable, optional*) – A function/transform that takes in the target and transforms it.

`torchvision.datasets`

CelebA

CIFAR

Cityscapes

+ COCO

DatasetFolder

EMNIST

FakeData

Fashion-MNIST

Flickr

HMDB51

ImageFolder

ImageNet

Kinetics-400

KMNIST

LSUN

MNIST

Omniglot

PhotoTour

Places365

QMNI

SBD

SBU

STL10

SVHN

UCF101


USPS

VOC

MNIST

```
training_data = datasets.MNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)
test_data = datasets.MNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)

figure = plt.figure(figsize=(8, 8))
cols, rows = 3, 3
for i in range(1, cols * rows + 1):
    sample_idx = torch.randint(len(training_data),
                                size=(1,)).item()
    img, label = training_data[sample_idx]
    figure.add_subplot(rows, cols, i)
    plt.title(label)
    plt.axis("off")
    plt.imshow(img.squeeze(), cmap="gray")
plt.show()
```



Loop de Treino e Validação

```
loss_fn = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(params=lenet5.parameters(), lr=0.1)

torch.manual_seed(42)
train_time_start_on_cpu = timer()

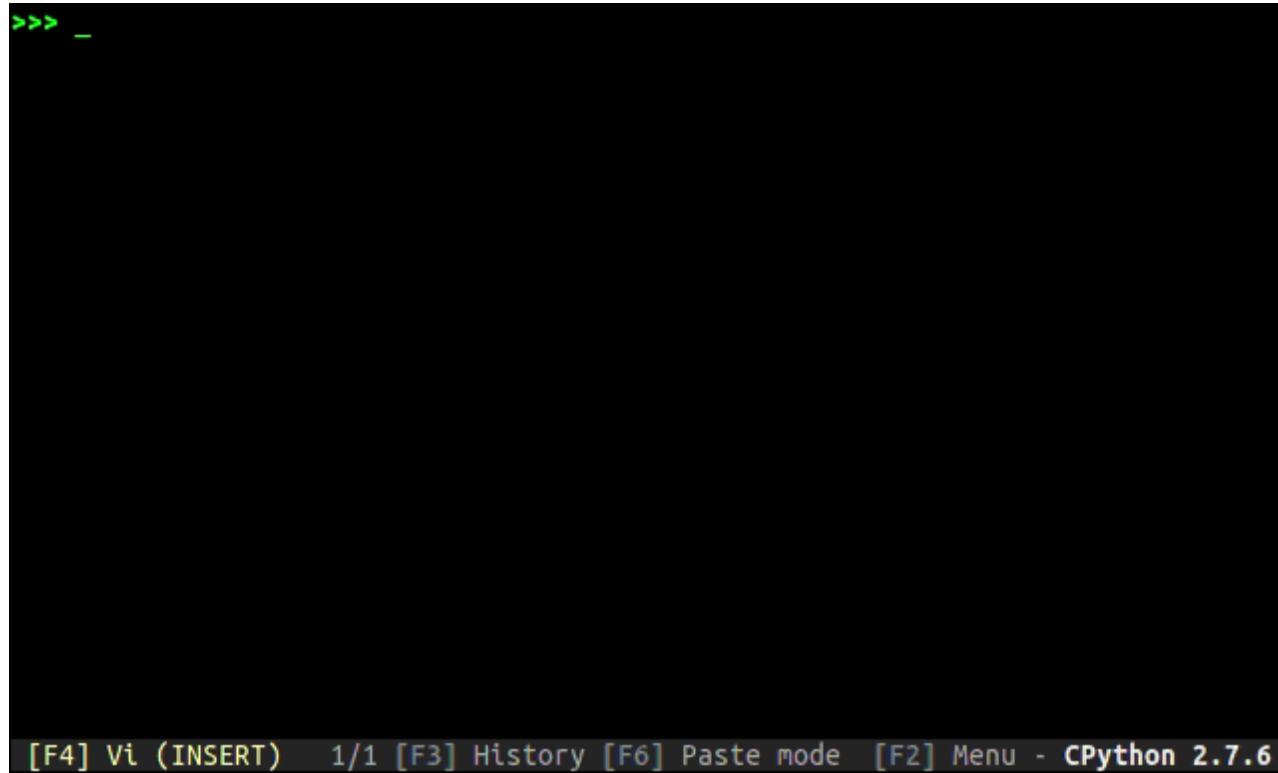
for epoch in range(EPOCHS):
    with tqdm(train_dataloader, desc=f'{epoch=}', unit='batch') as tqdm_epoch:
        train_loss = train_step(lenet5, tqdm_epoch, loss_fn, optimizer)
        test_loss, test_acc = test_step(lenet5, loss_fn, test_dataloader, device)

    print(f'Train loss: {train_loss:.5f} | Test loss: {test_loss:.5f}, Test acc: {test_acc*100:.4f}%')

train_time_end_on_cpu = timer()
total_train_time_model_0 = print_train_time(start=train_time_start_on_cpu,
                                             end=train_time_end_on_cpu,
                                             device=str(next(lenet5.parameters()).device))
```

tqdm

- É usual criar barras de processamento durante o treinamento e validação de redes neurais
- Um pacote python muito utilizado para isso é o tqdm



```
>>> _
```

[F4] Vl (INSERT) 1/1 [F3] History [F6] Paste mode [F2] Menu - CPython 2.7.6

Loop de Treino e Validação

```
def train_step(model, dataloader, loss_fn, optimizer, device: torch.device):
    train_loss = 0
    # Faz loop em todos os dados de treino
    for X, y in dataloader:
        X, y = X.to(device), y.to(device)
        model.train()
        # 1. Forward pass
        y_pred = model(X)
        # 2. Calcula loss por batch
        loss = loss_fn(y_pred, y)
        train_loss += loss
        # 3. zera gradientes anteriores
        optimizer.zero_grad()
        # 4. Backward Pass
        loss.backward()
        # 5. Otimizacao
        optimizer.step()
    train_loss /= len(dataloader)
    return train_loss
```

Loop de Treino e Validação

```
def test_step(model:torch.nn.Module, loss_fn:torch.nn.Module, dataloader:torch.utils.data.DataLoader,
device:torch.device):
    test_loss, test_acc = 0, 0
    model.eval()
    with torch.inference_mode():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            # 1. Forward pass
            test_pred = model(X)

            # 2. Loss
            test_loss += loss_fn(test_pred, y) # accumulatively add up the loss per epoch

            # 3. Computa a acuracia
            test_acc += accuracy(target=y,
                                preds=torch.softmax(test_pred,dim=1),
                                task='multiclass',
                                num_classes=10)

    test_loss /= len(dataloader)
    test_acc /= len(dataloader)
    return test_loss, test_acc
```

Salvando os pesos da rede

- Após o processo de treinamento (ou em algum checkpoint relevante), podemos gravar os pesos (e/ou a arquitetura) da nossa rede
- Usamos `torch.save` e `torch.load` para fazer a **serialização** e a **desserialização** da rede neural
 - Usa o módulo pickle

Dessa forma salvamos apenas os pesos

```
torch.save(lenet5.state_dict(), 'lenet5.pth')
```

```
model1 = LeNet5()  
model1.load_state_dict(torch.load('lenet5.pth'))
```

Dessa forma salvamos pesos e arquitetura

```
torch.save(lenet5, 'lenet5_model.pth')
```

```
model2 = torch.load('lenet5_model.pth')
```

Jupyter Notebook vs Script

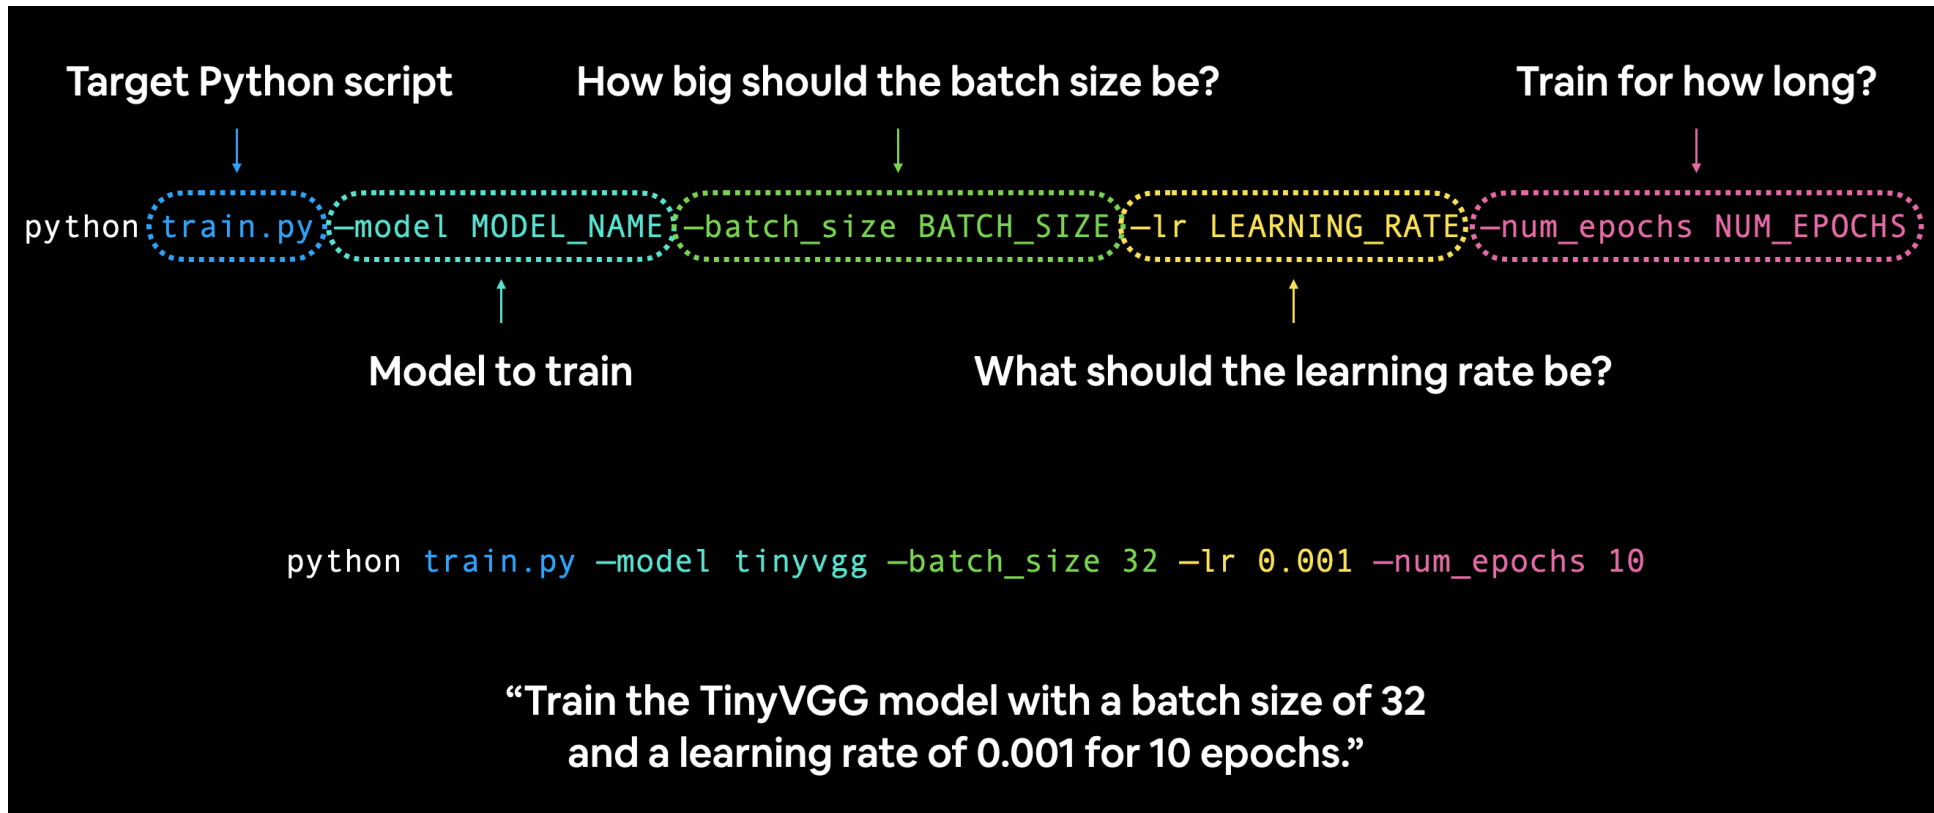
- Embora Jupyter Notebooks (*.ipynb) sejam mais interativos e fáceis para desenvolvimento, usualmente queremos a flexibilidade e facilidade de execução proporcionada por um script (*.py)
- É comum começarmos com jupyter e migrarmos para scripts

Jupyter Notebook vs Script

- Quando portamos nosso código para script precisamos tomar alguns cuidados
 - Modularizar o código separando em scripts diferentes cada uma das funcionalidades
 - `model.py`
 - `data_setup.py`
 - `train.py`
 - `eval.py`
 - Fornecer argumentos nos scripts para parâmetros que podem ser modificados em diferentes execuções do treinamento (`argparse` ou `typedargs` são soluções interessantes)
 - Evitar deixar códigos “solto” dentro do script

Jupyter Notebook vs Script

- Uma chamada para um script de treinamento de uma rede neural pode ser parecido com o seguinte



Jupyter Notebook vs Script

```
if __name__ == '__main__':  
    parser = argparse.ArgumentParser()  
    parser.add_argument('--batch_size', '-b', type=int, default=64,  
                        help='batch size for training')  
    parser.add_argument('--epochs', '-e', type=int, default=6,  
                        help='Training Epochs')  
    parser.add_argument('--num_workers', '-w', type=int, default=2,  
                        help='Number of workers for dataloader')  
    parser.add_argument('--force_cpu', '-c', type=bool, default=False,  
                        help='flag to force processing only in cpu')  
    args = parser.parse_args()  
  
    main(args.batchsize, args.epoch, args.workers, args.force_cpu)
```

```
(linux-torch) → script_example git:(main) X python main.py --workers 8|
```

torch.vision.models

- Temos disponível em torch.vision.models diversas arquiteturas e modelos pré-treinados
- É possível iniciar com um desses modelos e fazer modificações e algumas camadas para possibilitar modificações na tarefa realizada
 - Ex: modificar o número de classes do classificador no final da rede

- AlexNet
- ConvNeXt
- DenseNet
- EfficientNet
- EfficientNetV2
- GoogLeNet
- Inception V3
- MaxVit
- MNASNet
- MobileNet V2
- MobileNet V3
- RegNet
- ResNet
- ResNeXt
- ShuffleNet V2
- SqueezeNet
- SwinTransformer
- VGG
- VisionTransformer
- Wide ResNet

Algumas ferramentas úteis

- Para uso em terminal
 - **tmux** para permitir execução de diversos terminais simultaneamente e evitar que um problema na conexão finalize um processamento
 - **nvidia-smi** para monitorar as GPUs
 - **gpustat** pacote python para monitorar as GPUs (usa o nvidia-smi)
 - **htop** para avaliar uso de processadores e memória

- <https://github.com/tmux/tmux/wiki>
- <https://github.com/wookayin/gpustat>
- <https://htop.dev/>

The screenshot shows a terminal window with a tmux session. The top part shows the output of a python script (teste.py) running in a tmux window. The script outputs the device (cuda), the length of the train and test dataloaders, and the progress of an epoch (46%). The bottom part shows the output of the htop command, displaying system statistics and a list of running processes.

```
(linux-torch) → 02_cnn_dl git:(main) X python teste.py
device='cuda'
Length of train dataloader: 938 batches of 64
Length of test dataloader: 157 batches of 64
epoch=0: 46% |██████████| 429/938 [00:08<00:03, 145.16batch/s]
```

```
note-lsk Tue Mar 19 17:25:37 2024 532.09
[0] NVIDIA GeForce MX350 | 66°C, 28 % | 768 / 2048 MB | python/16139(?M)
```

```
1 [|||||] 39.4% 5 [|||||] 34.0%
2 [|||||] 35.3% 6 [|||||] 29.5%
3 [|||||] 36.7% 7 [|||||] 33.3%
4 [|||||] 17.2% 8 [|||||] 35.6%
Mem[|||||] 7.05G/7.63G Tasks: 22, 33 thr; 5 running
Swp[|||||] 2.00G/2.00G Load average: 1.45 1.02 0.60
Uptime: 12:44:56
```

| PID | USER | PRI | NI | VIRT | RES | SHR | S | CPU% | MEM% | TIME+ | Command |
|-------|-------|-----|----|-------|-------|-------|---|------|------|---------|-----------------|
| 16139 | lucas | 20 | 0 | 28.0G | 3482M | 412M | R | 104. | 44.6 | 0:09.23 | python teste.py |
| 16241 | lucas | 20 | 0 | 28.0G | 3482M | 412M | S | 33.6 | 44.6 | 0:00.71 | python teste.py |
| 16195 | lucas | 20 | 0 | 28.0G | 3482M | 412M | R | 24.2 | 44.6 | 0:00.40 | python teste.py |
| 16162 | lucas | 20 | 0 | 25.2G | 1818M | 13896 | S | 23.5 | 23.3 | 0:00.45 | python teste.py |
| 16178 | lucas | 20 | 0 | 25.2G | 1816M | 12268 | S | 20.8 | 23.3 | 0:00.46 | python teste.py |

```
F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice F8Nice +F9Kill F10Quit
[0] 0:python* "note-lsk" 17:25 19-Mar-24
```

Referências

- Documentação Padrão do PyTorch
 - <https://pytorch.org/docs/stable/generated/>
- PyTorch Fundamentals
 - https://www.learnpytorch.io/00_pytorch_fundamentals/
- tqdm
 - <https://tqdm.github.io/>
 - <https://towardsdatascience.com/training-models-with-a-progress-a-bar-2b664de3e13e>